

experimenting2_intervention

April 3, 2016

```
In [76]: # Project 2: Supervised Learning
        ### Building a Student Intervention System
```

0.1 1. Classification vs Regression

Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?

```
In [77]: # Import libraries
import numpy as np
import pandas as pd
```

```
In [78]: # Read student data
student_data = pd.read_csv("student-data.csv")
print "Student data read successfully!"
# Note: The last column 'passed' is the target/label, all other are feature columns
```

Student data read successfully!

Now, can you find out the following facts about the dataset? - Total number of students - Number of students who passed - Number of students who failed - Graduation rate of the class (%) - Number of features
Use the code block below to compute these values. Instructions/steps are marked using **TODOs**.

```
In [79]: # TODO: Compute desired values - replace each '?' with an appropriate expression/function call

n_students = len(student_data.index)
n_features = len(student_data.columns)
n_passed = sum([1 for y in student_data['passed'] if y == 'yes'])
n_failed = sum([1 for n in student_data['passed'] if n == 'no'])
grad_rate = 100.*n_passed/(n_passed + n_failed)

print "Total number of students: {}".format(n_students)
print "Number of students who passed: {}".format(n_passed)
print "Number of students who failed: {}".format(n_failed)
print "Number of features: {}".format(n_features)
print "Graduation rate of the class: {:.2f}%".format(grad_rate)
```

```
Total number of students: 395
Number of students who passed: 265
Number of students who failed: 130
Number of features: 31
Graduation rate of the class: 67.09%
```

0.2 3. Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

0.2.1 Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Let's first separate our data into feature and target columns, and see if any features are non-numeric.

Note: For this dataset, the last column ('passed') is the target or label we are trying to predict.

```
In [80]: # Extract feature (X) and target (y) columns
feature_cols = list(student_data.columns[:-1]) # all columns but last are features
target_col = student_data.columns[-1] # last column is the target/label
print "Feature column(s):-\n{}".format(feature_cols)
print "Target column: {}".format(target_col)

X_all = student_data[feature_cols] # feature values for all students
y_all = student_data[target_col] # corresponding targets/labels
print "\nFeature values:-"
print X_all.head() # print the first 5 rows
```

Feature column(s):-

['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'Mjob', 'Fjob', 'reason', 'guar', 'schstipend', 'passed']

Target column: passed

Feature values:-

| | school | sex | age | address | famsize | Pstatus | Medu | Fedu | Mjob | Fjob | \ |
|---|--------|-----|-----|---------|---------|---------|------|------|---------|----------|---|
| 0 | GP | F | 18 | U | GT3 | A | 4 | 4 | at_home | teacher | |
| 1 | GP | F | 17 | U | GT3 | T | 1 | 1 | at_home | other | |
| 2 | GP | F | 15 | U | LE3 | T | 1 | 1 | at_home | other | |
| 3 | GP | F | 15 | U | GT3 | T | 4 | 2 | health | services | |
| 4 | GP | F | 16 | U | GT3 | T | 3 | 3 | other | other | |

| | ... | higher | internet | romantic | famrel | freetime | goout | Dalc | Walc | health | \ |
|---|-----|--------|----------|----------|--------|----------|-------|------|------|--------|---|
| 0 | ... | yes | no | no | 4 | 3 | 4 | 1 | 1 | 3 | |
| 1 | ... | yes | yes | no | 5 | 3 | 3 | 1 | 1 | 3 | |
| 2 | ... | yes | yes | no | 4 | 3 | 2 | 2 | 3 | 3 | |
| 3 | ... | yes | yes | yes | 3 | 2 | 2 | 1 | 1 | 5 | |
| 4 | ... | yes | no | no | 4 | 3 | 2 | 1 | 2 | 5 | |

absences

| | |
|---|----|
| 0 | 6 |
| 1 | 4 |
| 2 | 10 |
| 3 | 2 |
| 4 | 4 |

[5 rows x 30 columns]

0.2.2 Preprocess feature columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply yes/no, e.g. `internet`. These can be reasonably converted into 1/0 (binary) values.

Other columns, like `Mjob` and `Fjob`, have more than two values, and are known as categorical variables. The recommended way to handle such a column is to create as many columns as possible values (e.g. `Fjob_teacher`, `Fjob_other`, `Fjob_services`, etc.), and assign a 1 to one of them and 0 to all others.

These generated columns are sometimes called dummy variables, and we will use the `pandas.get_dummies()` function to perform this transformation.

- Given what you know about the data so far, why did you choose this model to apply?
- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.

Produce a table showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.

Note: You need to produce 3 such tables - one for each model.

In [83]: *# Helper Functions*

```
import time
from sklearn.metrics import f1_score

# Return the classifier's training time
def timeTraining(clf, X_train, y_train):
    start = time.time()
    clf.fit(X_train, y_train)
    end = time.time()
    return "{:.3f}".format(end - start)

# Return the classifier's predictions and prediction time
def predictAndTime(clf, features):
    start = time.time()
    y_pred = clf.predict(features)
    end = time.time()
    return y_pred, "{:.3f}".format(end - start)

# Return the f1 score for the target values and predictions
def F1(target, prediction):
    return f1_score(target.values, prediction, pos_label='yes')
```

In [84]: *# To get my bearings I wanted to try most of the classifiers seen in class out of the box.
I excluded neural net since its recommended use requires the features to be scaled*

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier

# Array of classifiers
clfs = [DecisionTreeClassifier(criterion = "entropy"),
        SVC(C = 1.0, kernel="rbf"),
        GaussianNB(),
        AdaBoostClassifier(),
        KNeighborsClassifier(n_neighbors = 3)]

#Gathering Table column and index labels
classifier_names = [clf.__class__.__name__ for clf in clfs]
benchmarks = ["Training time", "F1 score training set", "Prediction time", "F1 score test set"]
table = pd.DataFrame(columns = classifier_names, index = benchmarks)

# Fit Classifiers
for clf in clfs:
    classifier = clf.__class__.__name__
```

```

t_train    = timeTraining(clf, X_train, y_train)
pred_train_set = predictAndTime(clf, X_train)[0]
pred_test_set, t_test = predictAndTime(clf,X_test)

table[classifier]['Training time']      = t_train
table[classifier]['F1 score training set'] = F1(y_train, pred_train_set)
table[classifier]['Prediction time']     = t_test
table[classifier]['F1 score test set']   = F1(y_test, pred_test_set)

```

In [85]: `from IPython.display import display, HTML`

`display(table)`

`#HTML(table.to_html())`

#Out of the box, with no tuning, it seems hard to differentiate their performance. Aside from

#slower training times seen in SVC and Adaboost, they pretty much all give similar f1 test scores

#SVC and Adaboost slightly ahead of the rest.

| | DecisionTreeClassifier | SVC | GaussianNB | \ |
|-----------------------|------------------------|----------|------------|---|
| Training time | 0.006 | 0.016 | 0.002 | |
| F1 score training set | 1 | 0.869198 | 0.808824 | |
| Prediction time | 0.000 | 0.004 | 0.001 | |
| F1 score test set | 0.757576 | 0.758621 | 0.75 | |

| | AdaBoostClassifier | KNeighborsClassifier |
|-----------------------|--------------------|----------------------|
| Training time | 0.136 | 0.001 |
| F1 score training set | 0.868778 | 0.886878 |
| Prediction time | 0.006 | 0.002 |
| F1 score test set | 0.779412 | 0.721805 |

In [86]: *# Helper function makeTable*

```
def makeTable(clf, training_sizes, X_tr, X_t, y_tr, y_t):
```

```
    #Gathering column and row labels for the table
```

```
    benchmarks = ["Training time", "F1 score training set","Prediction time", "F1 score test set"]
```

```
    size_labels = ["Training samples: {}".format(s) for s in training_sizes]
```

```
    table = pd.DataFrame(columns = benchmarks, index = size_labels)
```

```
    for i, size in enumerate(training_sizes):
```

```
        #Use only the first "size" number of samples
```

```
        X_train, X_test, y_train, y_test = [df.iloc[:size] for df in [X_tr, X_t, y_tr, y_t]]
```

```
        #Compute benchmarks
```

```
        t_train    = timeTraining(clf, X_train, y_train)
```

```
        pred_train_set = predictAndTime(clf, X_train)[0]
```

```
        pred_test_set, t_test = predictAndTime(clf,X_test)
```

```
        #fill table
```

```
        table['Training time'][i]      = t_train
```

```
        table['F1 score training set'][i] = F1(y_train, pred_train_set)
```

```
        table['Prediction time'][i]     = t_test
```

```
        table['F1 score test set'][i]   = F1(y_test, pred_test_set)
```

```
    return table
```

In [87]: *# Test Classifiers with increasing data set size*

```
training_sizes = [50,100,150,200,250,300]
```

```

for clf in clfs:
    print clf.__class__.__name__
    table = makeTable(clf, training_sizes, X_train, X_test, y_train, y_test)
    display(table)

```

DecisionTreeClassifier

| | Training time | F1 score training set | Prediction time | \ |
|-----------------------|---------------|-----------------------|-----------------|---|
| Training samples: 50 | 0.002 | 1 | 0.001 | |
| Training samples: 100 | 0.003 | 1 | 0.000 | |
| Training samples: 150 | 0.004 | 1 | 0.000 | |
| Training samples: 200 | 0.004 | 1 | 0.000 | |
| Training samples: 250 | 0.005 | 1 | 0.000 | |
| Training samples: 300 | 0.005 | 1 | 0.000 | |

| | F1 score test set |
|-----------------------|-------------------|
| Training samples: 50 | 0.677966 |
| Training samples: 100 | 0.694915 |
| Training samples: 150 | 0.703125 |
| Training samples: 200 | 0.753846 |
| Training samples: 250 | 0.710744 |
| Training samples: 300 | 0.772727 |

SVC

| | Training time | F1 score training set | Prediction time | \ |
|-----------------------|---------------|-----------------------|-----------------|---|
| Training samples: 50 | 0.001 | 0.90625 | 0.001 | |
| Training samples: 100 | 0.003 | 0.85906 | 0.002 | |
| Training samples: 150 | 0.003 | 0.870813 | 0.002 | |
| Training samples: 200 | 0.004 | 0.869281 | 0.002 | |
| Training samples: 250 | 0.006 | 0.879177 | 0.002 | |
| Training samples: 300 | 0.008 | 0.869198 | 0.002 | |

| | F1 score test set |
|-----------------------|-------------------|
| Training samples: 50 | 0.738462 |
| Training samples: 100 | 0.783784 |
| Training samples: 150 | 0.771429 |
| Training samples: 200 | 0.77551 |
| Training samples: 250 | 0.758621 |
| Training samples: 300 | 0.758621 |

GaussianNB

| | Training time | F1 score training set | Prediction time | \ |
|-----------------------|---------------|-----------------------|-----------------|---|
| Training samples: 50 | 0.001 | 0.666667 | 0.000 | |
| Training samples: 100 | 0.001 | 0.854962 | 0.000 | |
| Training samples: 150 | 0.001 | 0.808743 | 0.000 | |
| Training samples: 200 | 0.001 | 0.832061 | 0.001 | |
| Training samples: 250 | 0.001 | 0.817647 | 0.000 | |
| Training samples: 300 | 0.001 | 0.808824 | 0.000 | |

| | F1 score test set |
|----------------------|-------------------|
| Training samples: 50 | 0.468085 |

| | |
|-----------------------|----------|
| Training samples: 100 | 0.748092 |
| Training samples: 150 | 0.736842 |
| Training samples: 200 | 0.713178 |
| Training samples: 250 | 0.746269 |
| Training samples: 300 | 0.75 |

AdaBoostClassifier

| | Training time | F1 score training set | Prediction time \ |
|-----------------------|---------------|-----------------------|-------------------|
| Training samples: 50 | 0.099 | 1 | 0.005 |
| Training samples: 100 | 0.093 | 0.953846 | 0.006 |
| Training samples: 150 | 0.097 | 0.912821 | 0.006 |
| Training samples: 200 | 0.117 | 0.882562 | 0.006 |
| Training samples: 250 | 0.153 | 0.886427 | 0.007 |
| Training samples: 300 | 0.158 | 0.868778 | 0.008 |

| | F1 score test set |
|-----------------------|-------------------|
| Training samples: 50 | 0.645161 |
| Training samples: 100 | 0.72 |
| Training samples: 150 | 0.757576 |
| Training samples: 200 | 0.805755 |
| Training samples: 250 | 0.776978 |
| Training samples: 300 | 0.779412 |

KNeighborsClassifier

| | Training time | F1 score training set | Prediction time \ |
|-----------------------|---------------|-----------------------|-------------------|
| Training samples: 50 | 0.001 | 0.8 | 0.001 |
| Training samples: 100 | 0.001 | 0.823529 | 0.002 |
| Training samples: 150 | 0.001 | 0.816327 | 0.002 |
| Training samples: 200 | 0.001 | 0.86121 | 0.003 |
| Training samples: 250 | 0.001 | 0.889503 | 0.003 |
| Training samples: 300 | 0.001 | 0.886878 | 0.003 |

| | F1 score test set |
|-----------------------|-------------------|
| Training samples: 50 | 0.761905 |
| Training samples: 100 | 0.666667 |
| Training samples: 150 | 0.677419 |
| Training samples: 200 | 0.666667 |
| Training samples: 250 | 0.711111 |
| Training samples: 300 | 0.721805 |

In [88]: # 3 chosen classifiers: DecisionTree, SVC, and AdaBoost

The models were relatively stable as training size increased. KNN had a strange u-shaped behavior with high score decreasing, and then rising again. The SVC seemed the most stable across different training data set sizes.

chosen_clfs = [KNeighborsClassifier()]

In [89]: # tree.export_graphviz(clfs[0].fit(X_train,y_train), out_file='tree.dot')

0.4 5. Choosing the Best Model

- Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?
- In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a Decision Tree or Support Vector Machine, how does it make a prediction).
- Fine-tune the model. Use Gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.
- What is the model's final F1 score?

In [90]: *# TODO: Fine-tune your model and report the best F1 score*

```
from sklearn import grid_search
from sklearn.metrics import make_scorer
```

```
r = np.arange
scorer = make_scorer(F1)
```

```
neigh_param = {'n_neighbors' : [10,20,25,30,40], 'weights' : ['uniform', 'distance'], 'p':[1,2
```

```
#Perform grid Search
```

```
def gridIt(clf, params):
    grid_clf = grid_search.GridSearchCV(clf, params, scorer)
    final_clf = grid_clf.fit(X_train, y_train).best_estimator_
    print final_clf
    y_pred, predict_t = predictAndTime(final_clf, X_test)
    print F1(y_test, y_pred), predict_t
    print '-----\n'
```

```
gridIt(chosen_clfs[0], neigh_param)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
    metric_params=None, n_jobs=1, n_neighbors=30, p=2,
    weights='uniform')
```

```
0.77027027027 0.003
```

```
-----
```

In []: